Stolpe,
Graal-LSP

Graal-LSP: A Zero-overhead Polyglot Language Server Implementation for the GraalVM

Graal-LSP: Eine Overhead-freie Language Server Implementierung für die GraalVM

by

Daniel Stolpe

A thesis submitted to the
Hasso Plattner Institute
at the University of Potsdam, Germany
in partial fulfillment of the requirements for the degree of

Master of Science in IT-Systems Engineering

Supervisors

Prof. Dr. Robert Hirschfeld
Dr. John Doe

Software Architecture Group
Hasso Plattner Institute
University of Potsdam, Germany

December 26, 2018

Abstract

The english abstract.

Zusammenfassung

Die Zusammenfassung auf deutsch.

Acknowledgments

I owe everything to my cat.

Contents

1. Introduction **1**
1.1. Contributions . . . 1
1.2. Structure of this Thesis . . . 2

2. Context **3**
2.1. Terms and Concepts . . . 3
2.2. Truffle and GraalVM . . . 3
2.3. Language Server Protocol . . . 4

3. Approach **7**
3.1. A Language-agnostic Language Server . . . 7
3.1.1. Language Server Initialization . . . 8
3.1.2. Parsing Source Code . . . 8
3.1.3. Diagnostics . . . 9
3.1.4. Mapping Caret Position to AST Node . . . 9
3.1.5. Static Completion of Globals and Locals . . . 10
3.1.6. Completion of Locals Using Run-time Information . . . 12
3.1.7. Completion of Object Properties . . . 14
3.1.8. Signature Help . . . 16
3.1.9. Hover Information . . . 17
3.1.10. Document Symbols . . . 17
3.1.11. Workspace Symbols . . . 18
3.1.12. Document Highlight . . . 18
3.1.13. definition . . . 19
3.1.14. Finding All References . . . 19
3.1.15. Triggering Run-time Information Collection . . . 20
3.2. Architecture + Language Server Bootstrapping . . . 21

4. Implementation **23**
4.1. Language Server Implementation . . . 23
4.1.1. Architecture . . . 23
4.1.2. Using runtime information . . . 23
4.1.3. LSP features implementation overview . . . 23
4.1.4. LSP client - VSCode extension example . . . 25
4.2. Requirements on Truffle Language Implementations . . . 25

5. Experience Report **27**
5.1. Between IDE and Language Workbench . . . 27
5.2. Completion response time . . . 27

5.3. Limitations 27

6. Related Work **29**
6.1. Helvetia 29
6.2. Self 29
6.3. Smalltalk+Squeak 29
6.4. Squimera 29
6.5. Language Workbenches 29
6.6. Run-time type collection 29
6.7. CTags 29

7. Conclusion and Future Work **31**

A. First Unimportant stuff. **35**

List of Figures

2.1. Truffle languages hosted by the Truffle framework on top of the GraalVM 5
2.2. example for how a tool and a language server communicate during a routine editing session 6

3.1. lang specific server vs lang-agnostic server 7
3.2. A Visual Studio Code extension for the Python language showing the completions of globals and locals available in a function scope, together with reserved keywords. 10
3.3. A Visual Studio Code extension for the Python language showing all detected symbols in the current document. There are different icons for modules, classes, functions and variables. 18

List of Tables

3.1. Selected Language Server Protocol (LSP) features and their descriptions 8

List of Listings

1. An example Python code snippet showing two functions to explain scopes. 11
2. A Python code snippet showing an object property completion example. 15
3. A Python code snippet example to illustrate migrating cached run-time information. 20

List of Abbreviations

LSP	Language Server Protocol
IDE	Integrated Development Environment
AST	Abstract Syntax Tree
API	Application Programming Interface
VM	Virtual Machine
URI	Uniform Resource Identifier
VSCode	Visual Studio Code
TDD	Test-Driven Development

1. Introduction

- language development is challenging task
- what makes a good language: amongst performance, std lib, ecosystem and community, tools (for language developers and users) are important
- who provides the tools? language developers vs. language users
- normally tools are developed for a specific language and cannot be reused or need lots of adaptation
- In the Truffle framework, tooling (instrumentation) is part of the framework and therefore tools are implemented language-agnostic once
- Truffle has already some language-agnostic tools implemented (debugger, profiler, ...) and new ones can be easily built
- Truffle tooling support does not interfere with implementing high-performance languages
- Compared to a full IDE, some tools are missing
- Implementing LSP can provide further IDE-like features
- We propose an approach to build a generic LSP server on top of Truffle with:
  - minimal to zero language runtime performance overhead
  - minimal API changes to Truffle
  - minimal impl and maintenance overhead for language developers
- Benefits for the language developers: get IDE-like tools with minimal effort
- Benefits for language users: can choose from wide range of LSP clients + consistent dev-tools experience for different Truffle languages
- As proof of concept, we have implemented Graal-LSP- A Zero-overhead Polyglot Language Server Implementation for the GraalVM
- ?: focus on dynamically type languages? Or is this focus already included in Truffle (most impls are dynamically type langs)?

1.1. Contributions

- Contribution in detail:
  - An approach to build a generic LSP server on top of Truffle
  - An implementation of said approach: Graal-LSP
  - Proposals to extend the Truffle Framework with new interfaces needed to support more LSP features
  - An experience report how well existing Graal/Truffle-implemented languages work with Graal-LSP(with and without implementing the said interface proposals)

1.2. Structure of this Thesis

2. Context

In this chapter we explain terms, concepts and technologies helping to understand our approach in Chapter 3 and the implementation of Graal-LSP in Chapter 4.

2.1. Terms and Concepts

Abstract Syntax Tree (AST) Parsers for programming languages usually transform plain text source code into an intermediate representation which can be processed by compilers, interpreters or programming environment tools. An AST is such a typical representation which stores structural information of the source code implicitly in its tree structure. Therefore, concrete syntactic elements needed to write parsable code in plain text can be omitted. In general, inner nodes of the AST represent programming language operations and leaf nodes store the values needed for these operations. Depending on the use case, an AST can store additional information in its nodes. For example the lexical source section from which a node was derived. Programming environment tools can use these information to display any kind of errors, warnings or hints directly in the plain text source code. An AST interpreter evaluates an AST by traversing its tree structure and executing the operations defined in the nodes.

Scopes and Frames A Scope defines a region of a program in which a set of name bindings is valid, e.g. a local variable name. Most programming languages have a concept of a global (program wide) scope and local scopes defined by functions, blocks or modules. During execution of a program, frames store the actual values of the name bindings of visible scopes.

Static vs. Dynamic?

VM?

2.2. Truffle and GraalVM

Truffle [13] is a language implementation framework written in Java. The language developer uses the framework to write an AST interpreter, which is able to optimize itself during interpretation [16]. To achieve that, Truffle uses runtime type information to rewrite AST nodes. Callable elements of a language, e.g. functions or methods, are modeled as root nodes. They can be executed by providing a frame object. A node can have a source section attribute which encodes the range in the source code file it

relates to. How source code is parsed and the AST is structured leaves Truffle in the responsibility of the language implementation.

Truffle has built-in support for developing language-agnostic tools by providing an instrumentation Application Programming Interface (API) [10]. It is designed to work well with the optimization techniques for language implementations so that instrumentation can be done with nearly zero-overhead [11]. Instruments can be installed by putting it on the Java class path and then be registered and unregistered dynamically as listeners for different events. For example, an instrument can be notified when a new source or source section is loaded, a statement is going to be executed or an expression will return a value. It is also possible to alter the current execution flow and let a function return a different value when called under certain conditions. This enables building powerful tools like debuggers or profilers which need to be implemented only once and can then be used for all Truffle languages. To control which AST nodes of a language are available for instrumentation, language implementors need to mark nodes as instrumentable. A tagging mechanism allows adding meta information, like "this node represents a statement in my language", which can be used by instruments to filter for specific nodes. For example, a debugger can facilitate these information to implement stepping. The tools available at the date of writing are mostly runtime-tools like debuggers, dynamic profilers and language extensions [11].

Cross-language interoperability is another feature of the Truffle framework [3]. Language implementations can be hosted by the Truffle framework in a multi-language runtime and communicate via language-independent messages. Foreign language objects can be accessed directly without marshalling.

To get high-performance language implementations, Truffle uses the language-independent just-in-time compiler named Graal. It has its roots in the research project of the Maxine VM [12]. Among other optimizations, Graal applies partial evaluation to compile stable AST subtrees into optimized machine code [14, 2]. These optimizations do not stop at language boundaries, so that the access of foreign language objects can also be optimized.

cite or footnote?

GraalVM[1] is the umbrella term for a HotSpot VM which is bundled with the Truffle language implementation and instrumentation framework and the Graal compiler, building a high-performance polyglot VM for applications written in a Truffle-implemented language [15]. Figure 2.1 visualizes that.

2.3. Language Server Protocol

In 2016 a group of IT companies, including Microsoft, Red Hat, and Codenvy, started to standardize the Language Server Protocol (LSP) [5]. This protocol describes an approach to tackle the problem of language and tool developers, which need to repeatedly implement language support (e.g. code completion, jump-to-definition and syntax error reporting) for every new code editor, Integrated Development Environment (IDE) or development tool. The LSP defines a *language server* and a *language*

[1] https://www.graalvm.org/

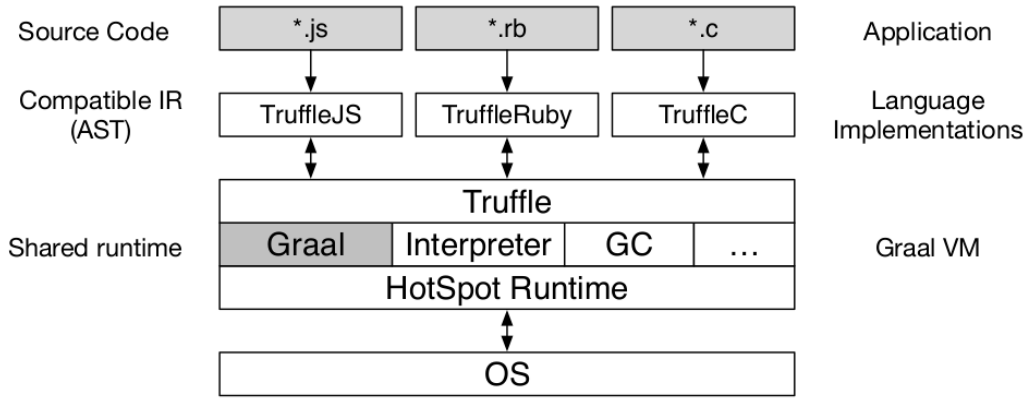


Figure 2.1.: Truffle languages hosted by the Truffle framework on top of the GraalVM

client. The server needs to be implemented only once for every language, providing the language-specific smarts. A development tool implementing a language client can make use of any language server and therefore support multiple languages with minimal effort.

At the date of writing, language servers for more than 50 languages exist[2], including almost all popular programming and scripting languages. Not all implementations support the full feature set defined in the LSP. Usually, these language server implementations support exactly one language, but the LSP allows multiple languages per server.

Language client and server communication LSP client and server communicate via JSON-RPC[3] requests, responses and notifications. The LSP version 3.7.0, which is the reference for our approach and implementation of Graal-LSP, specifies requests for typical features of an IDE: e.g. code and signature completion, navigational features (goto-definition, find all references and implementations, etc.), information on hover and code formatting. During initialization client and server can exchange which features are supported.

Figure 2.2 shows an exemplary communication between server and client. The client notifies the server that a file has been opened, the server parses the source code and pushes diagnostics to the client. Client-side changes to the source code are incrementally sent to the server.

LSP4J Our implementation of Graal-LSP uses LSP4J[4] version 0.4.x as Java binding for the Language Server Protocol. This framework maps the LSP requests and notifications to methods of Java interfaces so that the language server implementor does not have to deal with the JSON-RPC communication layer.

version

[2] https://microsoft.github.io/language-server-protocol/implementors/servers/ Sep. 2018
[3] https://www.jsonrpc.org/specification
[4] https://github.com/eclipse/lsp4j

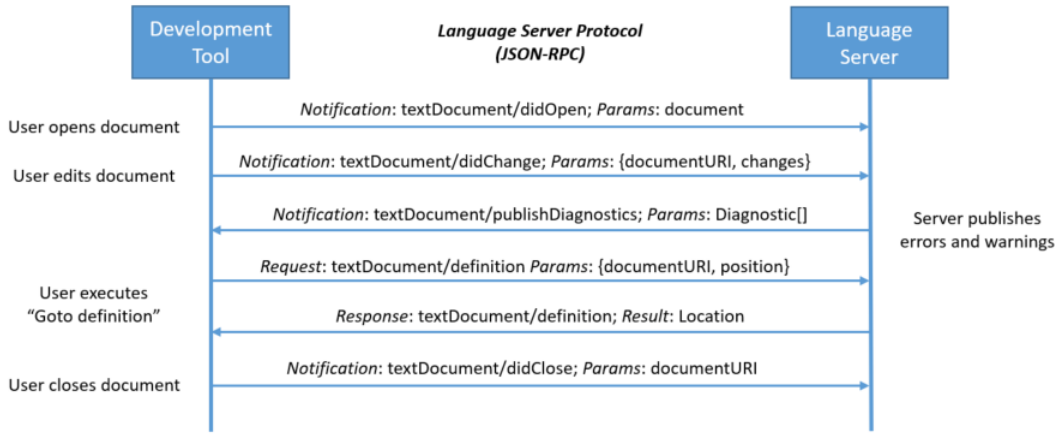


Figure 2.2.: example for how a tool and a language server communicate during a routine editing session

Visual Studio Code We used Microsoft's editor Visual Studio Code (VSCode)[5] to create an exemplary language client extension and experiment with Graal-LSP.

[5] https://code.visualstudio.com

3. Approach

- language implementation + tool development = challenging task
- we propose: implement LSP for Truffle
- typically one specific language server impl for one language
- Our approach: a language-agnostic language server impl for languages created by the language implementation framework Truffle, Figure 3.1

3.1. A Language-agnostic Language Server

- LSP enables arbitrary language clients to communicate with arbitrary language servers
- Normally a language server provides support for one specific language only
- To build a language-agnostic language server for Truffle we have to define what Truffle languages need to provide in general to support LSP
- LSP is trade-off - there are lang-specific IDEs with more/other features
- Therefore our solution will also be a trade-off using general language features available for all Truffle languages
- Truffle ASTs have runtime focus. Existing languages are implemented with performance focus. Infos not relevant for execution are dropped.

In the following sections we suggest how the existing Truffle language and instrumentation APIs can be facilitated to provide the needed data for LSP requests and notifications. When needed, we also suggest small API extensions to get access to additional language-specific information, which have not been exposed yet, due to fact that the existing tools are mostly run-time focused and do not require them.

Table 3.1 gives an overview over the most interesting requests and notifications of the LSP version 3.7.0 which we cover in our approach.

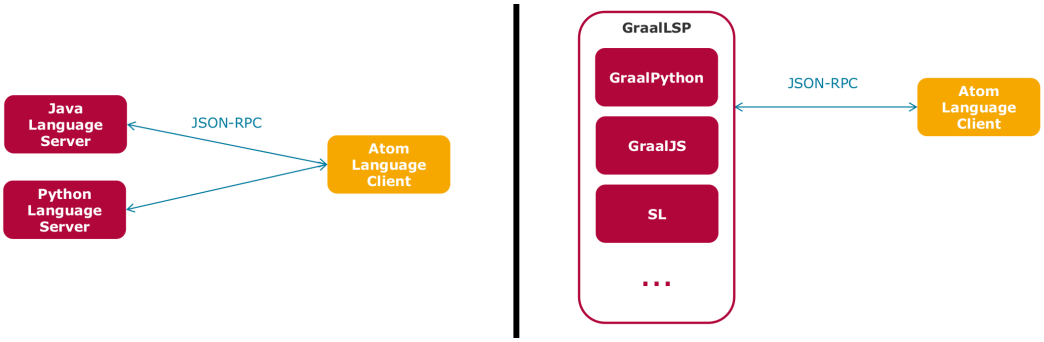


Figure 3.1.: lang specific server vs lang-agnostic server

Table 3.1.: Selected LSP features and their descriptions

Request/Notification	Description	Section
initialize	First request containing client capabilities. Response contains server capabilities.	3.1.1
didOpen	Notification sent from client to sever that a file was opened. Contains URI and content.	3.1.2
didChange	Notification sent from client to sever that a file has changed. Contains URI and changes.	3.1.2
diagnostics	Notification sent from server to client to display errors, warnings or hints.	3.1.3
completion	Completion request at caret position. Can be triggered automatically via registered completion trigger characters.	3.1.5
signatureHelp	Request to find signature information (e.g. parameter names and details for a function) at caret position. Is triggered via registered signatureHelp trigger characters.	3.1.8
hover	Request for information at caret-hovering position.	3.1.9
documentSymbol	Request for symbols of a specific source file (function names, variable names, module names, etc.).	3.1.10
symbol	Like documentSymbol, but for all files in the client workspace.	3.1.11
definition	Goto-definition request for functions, variables, etc.	3.1.13
references	Request to find all references of a symbol.	3.1.14
documentHighlight	Request to highlight all ranges related to a symbol (e.g. variable reads and writes).	3.1.12
codeLense	Request to provide executable actions at arbitrary source locations.	3.1.15

3.1.1. Language Server Initialization

The first request sent from the client to the server is the *initialize* request. Among other parameters, it contains client capabilities, telling the server which LSP features the client supports. The server responds to this by sending its own capabilities. When a server supports the *completion* and *signatureHelp* requests, the response needs to include trigger characters[1] which tell the client when it shall send a *completion* or *signatureHelp* request automatically. A language-specific language server implementation can provide these characters statically. In our language-agnostic approach, where Truffle languages can be installed dynamically, we have to collect these characters for all available languages during server initialization. Using the Truffle instrumentation environment, we can get a list of all registered languages, but an interface extension is needed to get the trigger characters for *completion* and *signatureHelp* requests of a language.

3.1.2. Parsing Source Code

To parse source code in Truffle, we need to define a `Source` object, containing the language and the content of the file. The language can be derived from its

[1] A trigger character is actually a string and can consist of multiple character symbols.

supported media types. The "didOpen" or "didChange" requests contain an Uniform Resource Identifier (URI) pointing to the concerning source file, which we use to find the required media type. The "didOpen" request also contains the file content. The "didChange" request contains only incremental changes if configured during initialization. Therefore we have to continuously store the content and apply all incoming changes to keep track of the client-side file state.

The "initialize" request can contain a workspace URI. If present, we parse all files with Truffle language-associated media types in that folder and its subfolders at start-up. To do so, we need to collect all media types of all installed languages via the Truffle environment.

Truffle's parsing result is a `CallTarget`, ready to be executed. It holds a reference to a root node of a runtime-focused AST.

In case of a parsing error, we can notify the client using the "diagnostics" notification.

3.1.3. Diagnostics

Truffle languages typically throw a `TruffleException` in case of a parsing error. The parsers are not designed to be syntax error tolerant and will stop parsing at the first syntax error. A `TruffleException` can have a reference to the location where the exception occurred in the AST. If available, we map these information to line and column of the original source code and send a *diagnostics* notification to the client. If the location is not available, we use the first column in the first line as reference for the notification. The error message needed for the notification is also derived from the `TruffleException`.

The *diagnostics* notification is also used to display code coverage information, which will be described in xyz. ref

3.1.4. Mapping Caret Position to AST Node

Most LSP requests are context-sensitive and include line and column as reference in the request. The language server has the responsibility to map these information to data structures of the language. For Truffle, this means mapping line and column to AST nodes via their associated source section. The challenge for our language-agnostic language server is that the Truffle framework leaves the creation of the runtime-focused ASTs completely to the language implementation. We cannot simply traverse an AST by starting at the `CallTarget`'s executable root node which we get as parsing result. There might be more root nodes, which will be executed dynamically at runtime, depending on the language implementation. However, we can use the Truffle instrumentation environment to access all already loaded nodes which have a source section, under the condition, that the source code is parsable.

Depending on the use cases described in the next sections, we are either interested in a node directly at the caret position ("hover", "definition" and other requests) or a node near the caret position ("completion" request for globals and locals). To find these nodes in a language-agnostic way, we have to define two assumptions.

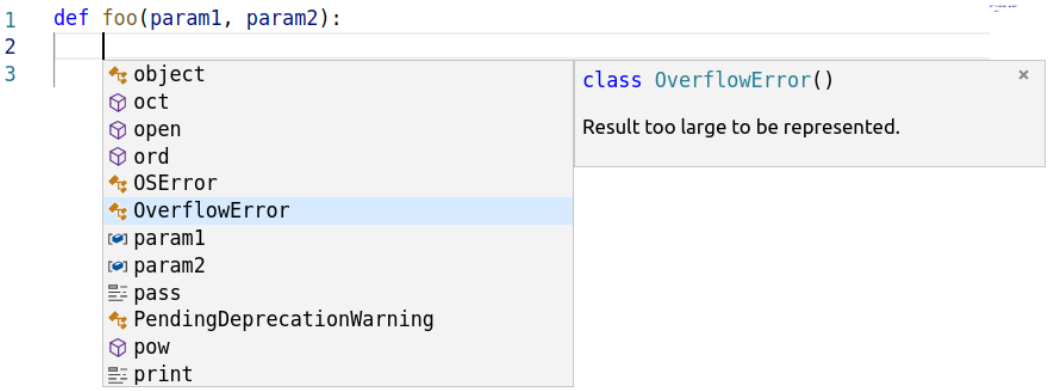


Figure 3.2.: A Visual Studio Code extension for the Python language showing the completions of globals and locals available in a function scope, together with reserved keywords.

Eager Source Code Parsing First of all, we assume, that parsing a Truffle `Source` object is done when we call the corresponding API and not lazily when the `CallTarget` is executed. Otherwise no nodes and source sections will be available via the Truffle instrumentation environment.

Consistent Source Section Hierarchy We assume furthermore, that source sections of ASTs follow the rule: if a node has a source section and is a child of another node, then the source section of the child is fully contained in all source sections of its direct parent or indirect parents. A parent without a source section in this hierarchy is skipped.

3.1.5. Static Completion of Globals and Locals

ref?

A context-sensitive code completion is one of the most important features to increase productivity of programmers . It suggests valid symbols regarding the current scope and thereby supports programmers to explore, for example, the built-in functions of a language, read their documentation or select predefined code snippets.

For Truffle, we have to distinguish between two kinds of code completion, both depending on the current caret position. The first one we call *completion of globals and locals*, described in this section. The second one we call *completion of object properties*, described in Section 3.1.7.

Completion of globals and locals is the completion you typically get in IDEs when the caret is placed in an empty line and the user presses a keyboard shortcut, e.g. `Ctrl+Space`. With *globals* we refer to all the symbols which we can derive from global scopes of languages and with *locals* we analogously refer to all symbols which we can derive from local scopes of languages. In the Python programming language [9] globals are, for example, built-in functions or classes. Locals of a function scope are function parameters or local variables. Locals of a class scope are fields or methods. Figure 3.2 exemplarily shows a code completion for globals and locals of a Python file in Visual Studio Code.

If the assumptions in Section 3.1.4 (eager source code parsing and a consistent source section hierarchy) are true, we can provide *static* code completion for globals and locals by using existing Truffle APIs. With *static*, we refer to a code completion, where we can only access lexical scopes to find locals. An approach to use runtime information for code completion, e.g. via a frame, is explained in Section 3.1.6.

Statically Deriving Local Scopes From a Node Global scopes are always accessible via the Truffle instrumentation environment. To find the relevant local scopes statically, we need to find a node which has a source section *relevantly near* the client's caret position. In the following we use some Python code snippets to explain what we mean by *relevantly near*.

Listing 1: An example Python code snippet showing two functions to explain scopes.

```
1  def foo(param1, param2):
2      x = 42
3
4      y = param1 + param2
5
6      print(x * y)
7
8
9  def bar():
10     return 42
```

listing 1 shows Python source code with a function definition `foo`, containing two parameters and two local variables in its scope and a function `bar` without definition of locals. Imagine that the caret is placed in the empty line 3 between the two variable assignments and code completion is triggered. The caret is clearly placed inside the scope of the function `foo` and we would assume that the code completion suggests, among globals, the two parameters `param1` and `param2`, and the local variable `x`. The local variable `y` is also in the function scope of `foo`, but is not written until line 4 and should therefore not be included in the suggestions. Whereas, if the caret is placed in line 5, `y` should be included.

label counter wrrong?

As already mentioned, Truffle cannot derive scopes directly from the caret position, but from a node. The difficulty is, that the scopes, which the Truffle API derives from a node, contain variables valid *at* that node. This means translated to our Python example, that if we use the node representing `x = 42` in line 2 to derive local scopes from, the scopes do not include the variable `x`. One could say, that we get all locals valid before the node's source section starts. Therefore we have to decide in a language-agnostic way, which node is near[2] the caret and provides the relevant information we need for the code completion.

Imagine code completion is triggered in line 3. Then the most inner node, which contains the caret, could be the function definition of `foo` ranging from line 1 to 6. The nearest node behind the caret could be the assignment of `y` in line 4 and the nearest node before the caret could be the assignment of `x` in line 2.

[2] For simplicity, we say in the following *a node near/after/before/containing the caret* instead of *a node which has a source section near/starting after/ending before/containing the caret position*

3. *Approach*

We suggest finding the node *relevantly near* the client's caret position following the rule:

1. If the caret is not contained in a node, we are not in a local scope. No completion of locals necessary. (Example: caret at line 8)
2. If contained, pick the node which is nearest behind the caret position and which is also an (indirect) child of the node which contains the caret. (Example: caret at line 3, picking node `y = param1 + param2`)
3. If not found, pick the node which is nearest before the caret. (Example: caret indented at line 7, picking node `y = param1 + param2`)
4. If still not found, use the node which contains the caret.

At this point, we have to note that this rule is only a best effort approach. Which locals are included in the lexical scopes provided via the Truffle API is depending on the language implementation. None the less, we are able to provide the needed information for a language-agnostic, static code completion of globals and locals with this approach.

Querying Properties of Scope Values For globals, we can not only access their names via the Truffle API, but retrieve additional information by querying the actual objects which are associated with these names in the global scopes. Based on the result, we can set the completion item kind supported by LSP, such as *variable*, *function*, or *class*.

A property, which is not supported by the current Truffle API, is documentation. Therefore, we propose to add a new Truffle API. For Python, this API could return the documentation string for functions, classes or modules.

To sum it up, with the proposed extension for documentation, we can support static code completion for locals without further details, but for globals with completion kind and documentation.

3.1.6. Completion of Locals Using Run-time Information

A detailed code completion for globals, locals, and properties of language elements is often implemented in IDEs using type information. For statically typed programming languages, like Java, type information can be easily derived from by parsing source code. The situation for dynamically typed languages, like Python, is different. The type of a Python variable is depending on its assigned value at run-time. There are language-specific approaches to infer types for dynamically typed languages , but these approaches do not match the language-agnostic requirements of the Truffle instrumentation API.

ref?

ref

In this section, we describe an approach on how to get and use run-time information to enhance code completion for locals. This approach will also build the basis for *completion of object properties* described in Section 3.1.7 and further requests described in later sections.

Approach and Challenges The basic idea to get access to run-time information (e.g. a frame) at caret position is, that we trigger source code execution while using Truffle instrumentation to listen for execution events occurring at a source section near the caret position. This approach implies some preconditions and challenges. We have to think about unexpected and unwanted side-effects during execution, the possibility of a long running or non-terminating execution, and the fact that the caret position might not be reached by the execution flow.

Sandboxing the Execution Environment The Truffle framework has an embedding-API which allows developers to configure sandboxing for the Truffle execution environment. At the moment, it supports restriction of the file system only, but in future versions of Truffle, the restriction of network I/O shall also be supported. It is also possible to install our own file system proxy, which enables having fine grained control over what Truffle languages are allowed to access and what not. We could therefore install a copy-on-write file system, which copies each file which shall be changed to a sandbox and redirects further read and writes to that copy. For now, we simply use a read-only file system, which prevents unwanted file changes outside the Truffle execution environment, when executing source code in our language server.

is embedding-API the right term?

Long Running or Non-terminating Executions A user, who triggers code completion, becomes impatient if the completion request takes longer than xseconds to return a result. The execution time of arbitrary source code might take a long time or might not finish at all, for example due to an endless loop. Therefore we have to explicitly stop execution when the execution flow reaches the desired source section. As a fallback, a time-out should ensure an upper bound of execution time by interrupting the execution. In that case, no run-time information can be provided for enhanced code completion, but an endless execution can be avoided.

value + ref

Unreachable Source Sections Depending on the source code, there might be source sections which are unreachable by the normal execution flow of the current source code file. Imagine, the Python code defined in the already mentioned listing 1, is the source code of a python file. If we execute that file, the two function definitions of `foo` and `bar` will be executed, but their function body will not be executed. Their is no function call to `foo` or `bar` defined in that file. Therefore, listening for execution events of the Truffle instrumentation environment to enhance our code completion in lines 2 to 6, which is the function body of `foo`, will not work. In our opinion, this is the normal case for file based programming languages. You typically split your program's source code into multiple files, having a limited set of files, which define runnable entry points, which then import elements of the other files.

One of these entry points are unit tests. Programmers typically write them to ensure the desired behavior of their programs and avoid regression of bugs. Therefore, unit tests should cover most parts of the source code. For our approach of getting run-time information to enhance code completion is the instrumented execution of unit tests a promising way. Most unit tests have a short execution time and are supposed

to be free of unwanted side-effects – or at least have a clean-up routine included. We only have to provide programmers a way to tell our language server how to execute the tests. This could be done via a separate configuration file or via a special comment in the first or last lines of a source code file, which points to a runnable script. This idea is inspired by *type harvesting* [4], a practical approach to obtaining typing information in dynamic programming languages by exploiting unit tests.

Because unit tests are not always available, especially if you start writing a small script or do not follow Test-Driven Development (TDD) principles , we use the current file, in which the code completion was triggered, as fallback for a runnable entry point.

ref to Kent Beck - Extreme Programming Explained: Embrace Change?

Caching Run-time Information Executing a script or unit tests for every *completion* request might still be to expensive and hinder the writing flow of the programmer. Therefore we need to cache collected run-time information. We suggest to use a simple mapping from source sections to frame snapshots which we set up using Truffle instrumentation. We know that different execution branches might reach the same source section with a frame, which has not only stored different values of the same type, but also values of different types. But for simplicity and performance, we only cache the first frame available during execution for every source section.

suggest to?

Having such a data structure, we can use the approach described in Section 3.1.4 to find a node near the caret an check via its source section for collected frames. If there is one, we can derive local run-time scopes from it. If there is none, we have to assume that the source section was not covered and use local lexical scopes as fallback.

3.1.7. Completion of Object Properties

In the previous sections we described the completion of globals and locals using static and run-time information. This section explains an approach to *complete object properties* in a language-agnostic way. With *completion of object properties* we refer to the typical completion you get in IDEs when inserting a trigger character, like a dot in Python or Java, after a variable. This way, you get suggestions for named properties available on the object bound to that variable. What a property actually is, depends on the Truffle language implementation. For a Python or Java object, properties could be all fields and methods which are defined by the object's class and super-classes. For a non-object-oriented language, this kind of completion might not be supported or be limited to special language elements, which understand some kind of *property* semantic. listing 2 shows a Python code snippet consisting of a class definition `Natural` with a field `n` and a method `next`. In the last line, a code completion could be triggered by an IDE automatically after the dot character was inserted. We would assume, that such a completion suggests the field `n` and the method `next` at least.

Approach and Challenges Truffle provides an API to list all property names of a language object and read or write the actual values associated with these property

names. To use this API, we have to solve some challenges. First of all we have to detect that an object property completion shall be triggered instead of a completion of globals and locals. Second, we have to identify the language element, i.e. a Truffle AST node, for which the completion of properties is desired. Finally, we have to obtain the actual Truffle object associated with that node, which might be available only at run-time.

Detecting an Object Property Completion To distinguish a completion request for globals and locals from one for object properties, we can use the API extension already suggested in Section 3.1.1, which allows us to ask a Truffle language for its completion trigger characters. Because the *initialize* response contains the trigger characters of all installed languages, the client will automatically trigger completion requests for all these characters. Therefore we have to check every request, if the trigger character belongs to the current language. If not, we have to skip the request.

The completion trigger character is only included in the request if the completion was triggered automatically and the client supports that optional feature. Therefore we have to check manually. To do so, we assume, that when a completion request arrives, we can distinguish the two kinds of completion by comparing one or multiple characters directly in front of the caret with the registered trigger characters of the current language. If there is a match, we try an object property completion.

Finding the Desired Node To find the AST node for which the completion of object properties is desired, the source code needs to be parsable. We assume, that this is not always the case when a completion is triggered, because an inserted completion character might cause a syntax error. Therefore we suggest to always revert the last textual change which was sent from the client via the *didChange* notification and then try to parse the reverted source code.

Listing 2: A Python code snippet showing an object property completion example.

```
1  class Natural:
2      n = 2
3
4      def next(self):
5          r = self.n
6          self.n = self.n + 1
7          return r
8
9  natural = Natural()
10 natural.
```

As an example, we refer again to listing 2. The source code shown is no valid Python source code because of the last dot character. Removing this dot makes the code parsable again.

We know that this approach does not fit for all situations and languages, but it is a first best-effort approach which works for languages like Python or JavaScript.

Naming limitations etc here? Or only in the experience report?

Now that we have a first approach to get parsable source code again, we can use the approach described in Section 3.1.4 to find a node directly in front of the current caret position.

Obtaining the Actual Properties The next step we have to do is getting the actual object associated with that node, which might be available only at run-time, to finally query it for its properties. The approach is basically the same as for the completion of locals using run-time information, described in Section 3.1.6. But there are some optimizations we can do.

First of all we can optimize property completion for literals. We can test the node in front of the caret, if it represents a literal and then query it for the actual literal object. This literal object can be used to obtain properties, if the language supports it. In some languages, for example Python, Java or Javascript, object property completion is possible for string literals.

If the literal-approach returns no results, we can check if there are already cached run-time information available (see Section 3.1.6). These information can be used by the Truffle instrumentation environment to perform an in-line evaluation of the source code associated with the node in front of the caret. As a prerequisite, a Truffle language needs to support in-line parsing of source code. The evaluation result can then be queried for properties.

If using cached run-time information returns no results, we use the actual approach described in Section 3.1.6, i.e. we execute the current script or unit tests and listen for execution events concerning the source section of the node in front of our caret. Translated to our Python example in listing 2, where we assume that a completion request was triggered by inserting a dot character in the last line, we are interested in the return value of the expression `natural` in front of that dot. Therefore we listen for execution events which are associated with the source section of `natural`. The resulting object can be queried for properties.

If there are still no results, because the source section might not be reachable for example, we can use a last fallback. We try to evaluate the source snippet associated with the node in front of the caret in the global scope. This will at least provide some property completion for globals and context-independent expressions.

name problems? context-depedent code might be evaluated in global scope yielding wrong results

To sum it up, with the approach described in this section, we are able to provide a basic, language-agnostic code completion for object properties. This approach includes reverting source code to a parsable state, executing that source code in a sandbox to listen for specific execution events using Truffle instrumentation, and querying the obtained run-time language object for its properties.

3.1.8. Signature Help

The *signatureHelp* request is sent from client to server to obtain signature information (e.g. function parameter names) at the caret position. The trigger characters for this request need to be defined in the response of the *initialize* request. Therefore, we have related challenges like the *completion* request for object properties. First of all, we have to distinguish if the trigger character, which caused the request, is belonging to the current language. If not, we can skip the request. Second, the sole insertion of the trigger character might cause a syntax error. Third, we have to obtain the actual run-time object to query it for signature information. These challenges can be analogously approached like we described in Section 3.1.6 for the completion of

object properties (bringing the source code into a parsable state, finding the node at the caret position and executing the source code until we have found the desired run-time object).

Unfortunately, there is no Truffle API yet, which provides signature information for an object. Therefore, we suggest to add a new Truffle API. Furthermore, we assume, that only callable elements of a language can have a signature. To mark these elements of a language, Truffle provides a `CallTag` for AST nodes. We can use this tag as a filter criterion for our search for the node at the caret position.

3.1.9. Hover Information

The *hover* request is sent from client to server to obtain information at the current (mouse) caret position. This request has the same prerequisite as the other context-depending requests. The source code needs to be parsable to find a node at the caret position. Having a node, we can check if the node's source section is associated with run-time information (see Section 3.1.6). If so, we can use these information (e.g. a frame) to build the *hover* response. If the hovered node represents a variable, we can use the frame to obtain the assigned value and include the string representations of that value and its meta-object in the response. Otherwise, we try to perform an in-line evaluation of the source code associated with the node. If the result is an object, which supports the suggested API to obtain signature information (see Section 3.1.8) or suggested API to obtain documentation (see Section 3.1.5), we can include signature and/or documentation, for example for a Python function, in the *hover* response.

For language developers or interested language users, we suggest to provide static hover information concerning the AST node at the caret position if no run-time information are available. This feature, which can be activated via an option, can be used to inspect the AST structure of a language and learn about which language elements in the source code are mapped to which kind of AST nodes. We used this feature extensively ourselves while implementing our approach.

is this remark fitting here?

3.1.10. Document Symbols

The *documentSymbol* request is sent from client to server to get all symbols available in a specific source file. An URI parameter is included in the request. Each symbol in the response must be associated with a source section. If the symbol, for example, relates to a function definition, the range of the definition including the function body could be taken as the source section. Additionally, a symbol kind can be added for each symbol in the response. Figure 3.3 shows how the *documentSymbol* response is displayed in Visual Studio Code with active Python extension.

To implement the *documentSymbol* request for Truffle languages, we have to think about which kind of symbols we can support with the current Truffle API. With the instrumentation environment we can access all loaded nodes which have a source section. We can filter them based on their meta information to find all root nodes. Because a root node is executable, we assume that the name of the root node is a

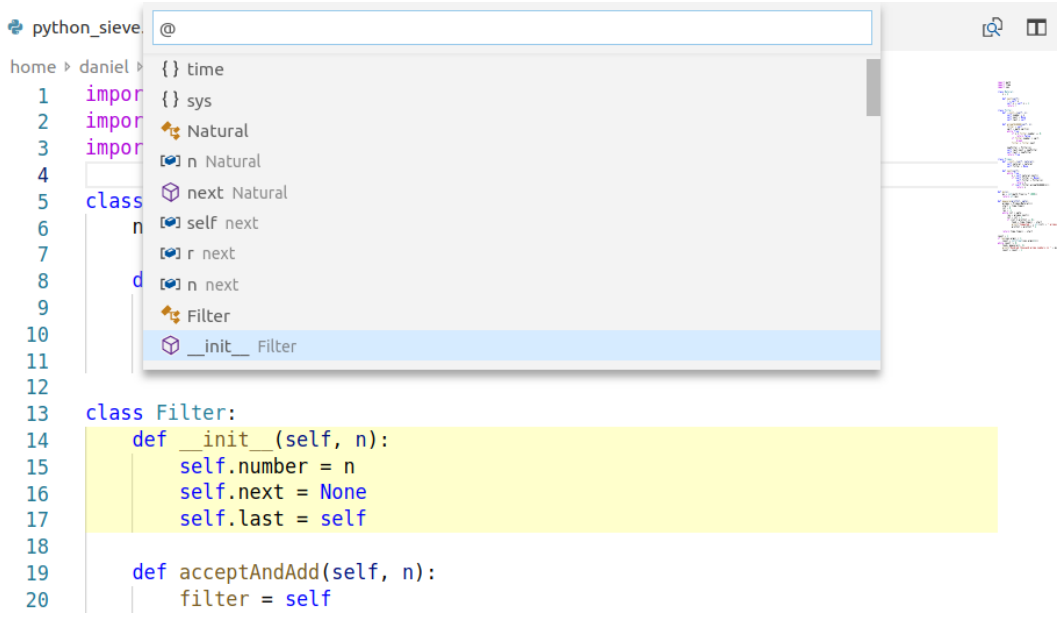


Figure 3.3.: A Visual Studio Code extension for the Python language showing all detected symbols in the current document. There are different icons for modules, classes, functions and variables.

valid symbol of a callable element of the corresponding Truffle language. However, to add a symbol kind we need more meta-information about the nodes.

Tagging Nodes with Meta-information To get the required meta-information, we suggest to add a new `Tag` for nodes introducing new symbols in a scope. There are already some tags existing to mark statements, expressions, calls, and root nodes. They are mainly used to implement debugger features like stepping. Using again the example of the Python language, the tag we suggest could be used to mark function and class definition nodes, first occurrences of a variable assignment node (Python has no variable declarations), or import nodes. This way we could support the *documentSymbol* request in a language-agnostic way with results equal to those language-specific ones already shown in Figure 3.3.

3.1.11. Workspace Symbols

The *symbol* request is the union of *documentSymbol* requests for all files in the workspace. Therefore, if we apply the approach explained in Section 3.1.10 for *documentSymbol*, we can easily add support for the workspace-wide *symbol* request.

3.1.12. Document Highlight

The *documentHighlight* request is sent from client to server to get a list of ranges, which shall be highlighted in the client. Additionally, we can define if the highlighted range is a read or write of a symbol. We want to support highlighting of variable

reads and writes for now. Unfortunately, there is no Truffle API yet to filter for nodes representing a read- or write-access of a variable. Therefore, we suggest to add two new tags to mark nodes which read or write a variable. These nodes also have to expose the identifier of the variable which they refer to. Having this, we can check if the node at the current caret position is tagged with one of these new tags. Then we can start a static search in the local scopes derived from that node for other tagged nodes with the same variable identifier. For every matching node, we add and response item, which includes the node's source range and the highlighting kind.

3.1.13. definition

- parsable code + find node at caret
- function vs variable
- function
  - dynamic approach: run-to-section to get function object
  - ask function object for its source section of definition
  - static fallback: search for symbols in workspace
- variable
  - not all languages have a variable definition syntax -> first write is definition
  - so we interpret goto-definition for variables as goto-first-write
  - find first write in scopes

3.1.14. Finding All References

The *references* request is sent from client to server to get all locations in the workspace, where the symbol at the caret position is also referenced. For now, we support a reference search for local variables and callable objects. As for many requests described before, the code needs to be parsable so that we can find a node a the current caret position.

For callable objects, we can search for nodes tagged with Truffle's `CallTag`. If there is such a tagged node at the caret, we can obtain its symbol and search all already loaded sources for other tagged nodes with the same symbol. Then we return all source section ranges of the matching nodes in the response.

For local variables, we need the tags suggested in Section 3.1.12 to search for nodes which represent a read- or write-access to a variable. If there is a node tagged with one of these two tags at the current caret position, we can search for all other variable references as described in Section 3.1.12 for the highlighting of variable reads and writes.

3.1.15. Triggering Run-time Information Collection

In Section 3.1.6 and the following ones, we described how to use cached run-time information for code completion, hover details and other requests. The question is, when to initialize and renew this cache. Therefore, we use the *codeLense* request, sent from client to server, to provide an interface for LSP-client users, so that they can trigger the renewal of the cache when they considered it reasonable. This could be the case, when the hover request returns incomplete or no results anymore.

The *codeLense* response includes a list of commands associated with display names and source ranges. It allows interspersing contextual actions directly into the source code. In a Visual Studio Code extension, it is used for example to insert a clickable link above a TypeScript function, if it violates some predefined code metrics. Clicking the link shows the metrics. We use it to insert such a code lense in the first line of the code, which will trigger the execution of the current source code or another script to collect run-time information.

Migrating Cached Information One challenge, we have not discussed so far, is that the cached run-time information are bound to specific source sections of AST nodes. Every time a *didChange* notification arrives, one or more sections of the source code will change. Depending on the change, the cached run-time information might become invalid for some source sections, but others might still be useful.

Imagine, we have collected run-time information for the Python code snippet in listing 3. Now, we edit this source snippet by inserting some code in line 3, which triggers an object property completion request for variable `anObject`. The collected run-time information in lines 1 and 2 are still valid, so that we can use them to perform an in-line evaluation for `anObject` to obtain its properties (see Section 3.1.7).

Listing 3: A Python code snippet example to illustrate migrating cached run-time information.

```
1 x = someCalculation(42)
2 anObject = expensiveOperation(x)
3
4 doSomething(y)
```

Imagine, the next change inserts a new empty line between lines 1 and 2. Our mappings of source sections to run-time information for all lines greater 1 are all invalid now. Therefore we suggest a simple approach to migrate these information. In the current case where a new line was inserted, this would be to simply adjust the off-by-one mismatch in the line numbers of the source sections following after that line. This is possible, if we address a source section only by its start and end line numbers and column offsets. We assume, that most edits will affect only their current local scope and there only the current line, so that we can apply the following rules for our mappings.

1. If a source section's end line is smaller than the start line of the change, we do nothing.

2. If a source section's start line is greater than the end line of the change, we check if the change adds or removes lines. If so, we adjust start and end line of the mapping by that difference.
3. If the range of the change is fully included in a source section, we adjust only the end line of the mapping by the difference of added or removed lines.
4. In all other cases where start or end lines of mapped source section and change overlap, we assume that the mapping is not valid anymore.

This way, we can migrate cached run-time information mappings for simple changes. But we can only assume that the adjusted mappings are still valid. The user has to trigger a new collection run manually. For example, if changes alter the execution flow, names of variables or their types.

3.2. Architecture + Language Server Bootstrapping

- graphic showing launcher, lsp server, truffle api communication
- comparison to other language server impls (built on top of language)

4. Implementation

4.1. Language Server Implementation

- lsp4j - Java binding for LSP (see Context)

4.1.1. Architecture

- Truffle-API + Graal-SDK
- Launcher, Instrument, Server
- lsp4j adapters

4.1.2. Using runtime information

Implementation of:

- Observe execution via instrumentation
- Side-effects -> Truffle-sandboxing
- Nested Contexts
- Evaluation strategies:
  - Literals
  - Coverage data
    - storing frames (copies)
    - source location format
    - migrating coverage data on source edit
  - Run-to-section
  - Globals

4.1.3. LSP features implementation overview

Feature impl matrix:

- TextDocumentService interface
  - completion
    - locals+globals vs object property completion (via checking completion character)
    - locals+globals completion
      - find nearest node

4. *Implementation*

        · check coverage data
        · globals only fallback
        · Used Truffle-API: findLocalScopes() + findGlobalScopes() + (parseInline() for coverage) + ForeignAccess.sendIsExecutable() + ForeignAccess.sendIsInstantiable() + ForeignAccess.sendIsNull()
        · Used new Truffle-API: env.boxPrimitive() + GetSignature + GetDocumentation

    * object property completion

        · fixing source code (undo last edit + re-parse)
        · find nearest node
        · tryDifferentEvalStrategies() (literals -> coverage data -> run-to-section -> globals)
        · Used Truffle-API: findLocalScopes() + findGlobalScopes()

  – hover
  – signatureHelp
  – definition
  – typeDefinition
  – implementation
  – references
  – documentHighlight
  – documentSymbol
  – codeAction
  – codeLense
  – formatting
  – rangeFormatting
  – onTypeFormatting
  – rename
  – didOpen
  – didChange
  – didClose
  – didSave
  – willSave
  – willSaveWaitUntil
  – documentLink
  – documentLinkResolve
  – documentColor
  – colorPresentation

- WorkspaceService interface

  – executeCommand
  – symbol
  – didChangeConfiguration
  – didChangeWatchedFiles

- didChangeWorkspaceFolders

4.1.4. LSP client - VSCode extension example

4.2. Requirements on Truffle Language Implementations

- Implement existing API:
  - ExpressionTag + StatementTag
  - findTopScopes()
  - findLocalScopes()
  - findMetaObject()
  - findSourceLocation()
  - parseInline()
- Implement proposed/new API:
  - getCompletionTriggerCharacters()
  - getSignatureHelpCharacters() (TODO)
  - boxPrimitive()
  - GetSignature + GetDocumentation custom message
  - DefinitionTag + ReadVariableTag + WriteVariableTag

5. Experience Report

5.1. Between IDE and Language Workbench

Language Workbench focus/aim: Efficient language (DSL) development, performance is not main focus vs. Graal/Truffle aim: Efficient execution

Assumption: Language Workbenches use static information only, our Truffle-based solution uses dynamic type information

5.2. Completion response time

5.3. Limitations

own chapter? or in 4?

- VSCode: Need to define supported languages in extension
- VSCode: Syntax highlighting defined at client-side
- LSP: No incremental sending of (completion) results
- Server: Only TCP (connecting to already running server), no std-in/-out impl
- Truffle: Only first syntax error is reported. Lang parsers fail on first error

6. Related Work

6.1. Helvetia

[8]

6.2. Self

6.3. Smalltalk+Squeak

6.4. Squimera

[7] [6]

6.5. Language Workbenches

[1]

6.6. Run-time type collection

[4]

6.7. CTags

- tagging a language object in an index file
- tag-file supported by editors
- allows user to locate the tagged object via name (goto-definition, list symbols)
- for new languages you need to define a regex or parser to find language objects
- only static analysis, no scoped completion (everything is suggested)

SemanticDB

7. Conclusion and Future Work

Bibliography

[1] Sebastian Erdweg, Tijs Van Der Storm, Markus Völter, Meinte Boersma, Remi Bosman, William R Cook, Albert Gerritsen, Angelo Hulshout, Steven Kelly, Alex Loh, et al. "The State of the Art in Language Workbenches". In: *International Conference on Software Language Engineering*. Springer. Springer International Publishing, 2013, pp. 197–217. DOI: 10.1007/978-3-319-02654-1_11.

[2] Yoshihiko Futamura. "Partial evaluation of computation process–an approach to a compiler-compiler". In: *Higher-Order and Symbolic Computation* 12.4 (1999), pp. 381–391. DOI: 10.1023/A:1010095604496.

[3] Matthias Grimmer, Chris Seaton, Roland Schatz, Thomas Würthinger, and Hanspeter Mössenböck. "High-performance cross-language interoperability in a multi-language runtime". In: *ACM SIGPLAN Notices*. Vol. 51. 2. ACM. 2015, pp. 78–90. DOI: 10.1145/2936313.2816714.

[4] Michael Haupt, Michael Perscheid, and Robert Hirschfeld. "Type harvesting: a practical approach to obtaining typing information in dynamic programming languages". In: *Proceedings of the 2011 ACM Symposium on Applied Computing*. ACM. 2011, pp. 1282–1289. DOI: 10.1145/1982185.1982464.

[5] Microsoft. *Language Server Protocol*. https://microsoft.github.io/language-server-protocol/. Visited on 2018-09-24. 2017.

[6] Fabio Niephaus, Tim Felgentreff, Tobias Pape, and Robert Hirschfeld. "Squeak Makes a Good Python Debugger: Bringing Other Programming Languages Into Smalltalk's Tools". In: *Companion to the first International Conference on the Art, Science and Engineering of Programming*. ACM. 2017, p. 27. DOI: 10.1145/3079368.3079402.

[7] Fabio Niephaus, Tim Felgentreff, Tobias Pape, Robert Hirschfeld, and Marcel Taeumel. "Live Multi-language Development and Runtime Environments". In: *arXiv preprint arXiv:1803.10200* 2.3 (Mar. 2018). DOI: 10.22152/programming-journal.org/2018/2/8.

[8] Lukas Renggli, Tudor Gîrba, and Oscar Nierstrasz. "Embedding languages without breaking tools". In: *European Conference on Object-Oriented Programming*. Springer. 2010, pp. 380–404. DOI: 10.1007/978-3-642-14107-2_19.

[9] Guido van Rossum. *Python Tutorial*. Tech. rep. CS-R9526. Centrum voor Wiskunde en Informatica (CWI), 1995.

[10] Michael L Van De Vanter. "Building debuggers and other tools: we can have it all". In: *Proceedings of the 10th Workshop on Implementation, Compilation, Optimization of Object-Oriented Languages, Programs and Systems*. ACM. 2015, p. 2. DOI: 10.1145/2843915.2843917.

[11] Michael Van De Vanter, Chris Seaton, Michael Haupt, Christian Humer, and Thomas Würthinger. "Fast, Flexible, Polyglot Instrumentation Support for Debuggers and other Tools". In: *arXiv preprint arXiv:1803.10201* (2018). DOI: 10.22152/programming-journal.org/2018/2/14.

[12] Christian Wimmer, Michael Haupt, Michael L Van De Vanter, Mick Jordan, Laurent Daynès, and Douglas Simon. "Maxine: An approachable virtual machine for, and in, java". In: *ACM Transactions on Architecture and Code Optimization (TACO)* 9.4 (2013), p. 30. DOI: 10.1145/2400682.2400689.

[13] Christian Wimmer and Thomas Würthinger. "Truffle: a self-optimizing runtime system". In: *Proceedings of the 3rd annual conference on Systems, programming, and applications: software for humanity*. ACM. 2012, pp. 13–14. DOI: 10.1145/2384716.2384723.

[14] Thomas Würthinger, Christian Wimmer, Christian Humer, Andreas Wöß, Lukas Stadler, Chris Seaton, Gilles Duboscq, Doug Simon, and Matthias Grimmer. "Practical Partial Evaluation for High-performance Dynamic Language Runtimes". In: *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation - PLDI 2017*. Vol. 52. 6. ACM. 2017, pp. 662–676. DOI: 10.1145/3062341.3062381.

[15] Thomas Würthinger, Christian Wimmer, Andreas Wöß, Lukas Stadler, Gilles Duboscq, Christian Humer, Gregor Richards, Doug Simon, and Mario Wolczko. "One VM to rule them all". In: *Proceedings of the 2013 ACM international symposium on New ideas, new paradigms, and reflections on programming & software*. ACM. 2013, pp. 187–204. DOI: 10.1145/2509578.2509581.

[16] Thomas Würthinger, Andreas Wöß, Lukas Stadler, Gilles Duboscq, Doug Simon, and Christian Wimmer. "Self-optimizing AST interpreters". In: *ACM SIGPLAN Notices*. Vol. 48. 2. ACM. 2012, pp. 73–82. DOI: 10.1145/2384577.2384587.

Appendix A.

First Unimportant stuff.

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Eigenständigkeitserklärung

Hiermit versichere ich, dass ich die vorliegende Arbeit selbständig verfasst sowie keine anderen Quellen und Hilfsmittel als die angegebenen benutzt habe.

Berlin, den 26. Dezember 2018

Daniel Stolpe

Todo list

- cite or footnote? . . . 4
- version . . . 5
- ref . . . 9
- ref? . . . 10
- label counter wrrong? . . . 11
- ref? . . . 12
- ref . . . 12
- is embedding-API the right term? . . . 13
- value + ref . . . 13
- ref to Kent Beck - Extreme Programming Explained: Embrace Change? . . . 14
- suggest to? . . . 14
- Naming limitations etc here? Or only in the experience report? . . . 15
- name problems? context-depedent code might be evaluated in global scope yielding wrong results . . . 16
- is this remark fitting here? . . . 17
- own chapter? or in 4? . . . 27